

Scalable parallel implementation of a coupling scheme using libMesh

Thiago MILANETTO SCHLITTNER, Régis COTTEREAU

Laboratoire MSSMat, CNRS, CentraleSupélec, Université Paris-Saclay

SIAM - CSE17



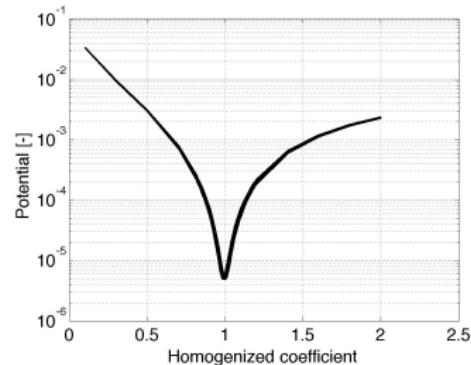
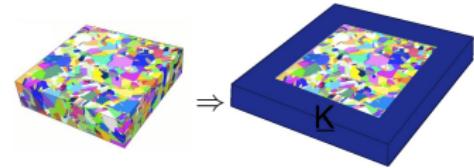
Funded by the French National Research Agency (projet number ANR-14-CE07-0007 CouEST)

Motivation

Homogenization

Exchange known heterogeneous model by unknown homogeneous model

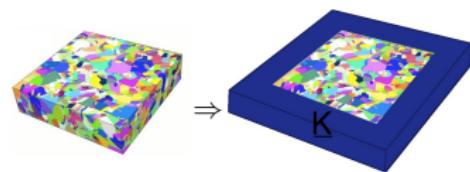
- Example:
 - **Micro:** heterogeneous, stochastic, non-linear model
 - **Macro:** homogeneous, deterministic model
→ homogenization of the micro model
- Models are linked through a coupling method



R. Cottreau, Int. J. Numer. Meth. Engng, 95, 1, 2013

Motivation

- Coupling method: take into account ...
 - different scales of the coupled models
 - non-conformal / incompatible meshes
- Latter is very important for stochastic models (random grain geometry realizations)



In this talk ...

- Present a scalable framework (**CArI**) to parallelize the coupling construction and to solve the coupled system

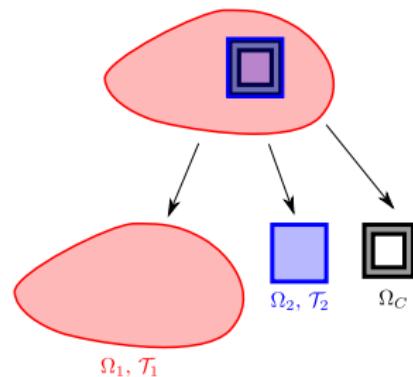
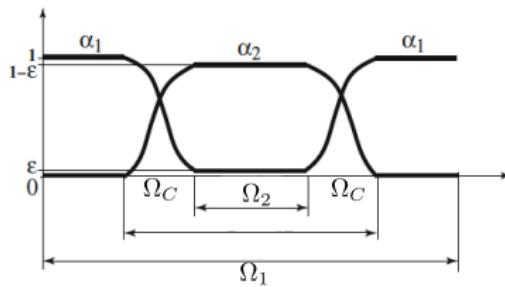
Outline

- Arlequin method: Description and formulation
- Intersection search: Parallelization and scaling
- Coupled solver: Implementation and scaling
- Example: linear heterogeneous anisotropic material
 - Same coupling formulation as non-linear model
 - Corresponds to one realization of a stochastic model

Other coupling methods have similar challenges (i.e. surface coupling)

Arlequin method

- Two overlapping models, defined on regions Ω_1, Ω_2
 - Meshes \mathcal{T}_1 and \mathcal{T}_2
- Coupling defined over region Ω_C
 - Volumetric coupling
 - Energy is partitioned between the models at the coupling region



D. Néron et al., Comput Mech, 2016

Arlequin method

- Formulation:

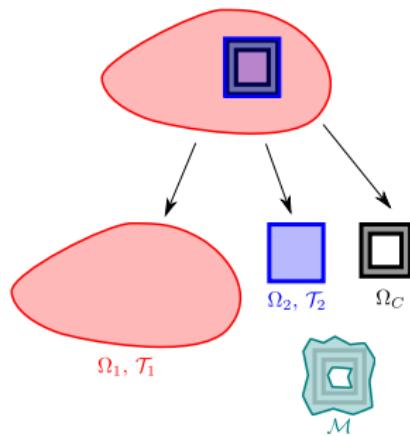
Find $(\mathbf{u}_1, \mathbf{u}_2, \Phi) \in \mathbf{V}_1 \times \mathbf{V}_2 \times \mathbf{M}$:

$$\begin{aligned} a_1(\mathbf{u}_1, \mathbf{v}) + c(\Phi, \mathbf{v}) &= f_1(\mathbf{v}) \quad \forall \mathbf{v} \in \mathbf{V}_1 \\ a_2(\mathbf{u}_2, \mathbf{v}) - c(\Phi, \mathbf{v}) &= f_2(\mathbf{v}) \quad \forall \mathbf{v} \in \mathbf{V}_2 \\ c(\psi, \mathbf{u}_1 - \mathbf{u}_2) &= 0 \quad \forall \psi \in \mathbf{M} \end{aligned}$$

- Coupling uses a *mediator space* \mathbf{M}

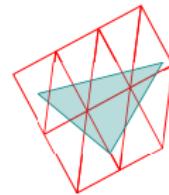
- mesh \mathcal{M}
- \mathbf{u}_2, ψ : stochastic models
- Coupling operator

$$c(\psi, \mathbf{v}) = \int_{\Omega^C} \kappa \left(\boldsymbol{\varepsilon}(\psi) : \boldsymbol{\varepsilon}(\mathbf{v}) + \frac{1}{e^2} \boldsymbol{\psi} \cdot \mathbf{v} \right) d\Omega$$



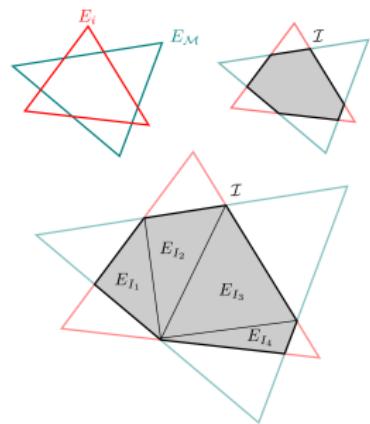
Arlequin method

$$c(\psi, \mathbf{v}) = \int_{\Omega^C} \kappa \left(\boldsymbol{\varepsilon}(\psi) : \boldsymbol{\varepsilon}(\mathbf{v}) + \frac{1}{e^2} \psi \cdot \mathbf{v} \right) d\Omega$$



- Spatial parts of ψ and \mathbf{v} are from different meshes! Cannot use quadrature from either!
- Solution: intersection mesh between \mathcal{T}_i and \mathcal{M} , **inside** Ω^C
- Only used to compute the integral

How to search and parallelize the intersections efficiently over hundreds of processors?



Intersection search: parallelization

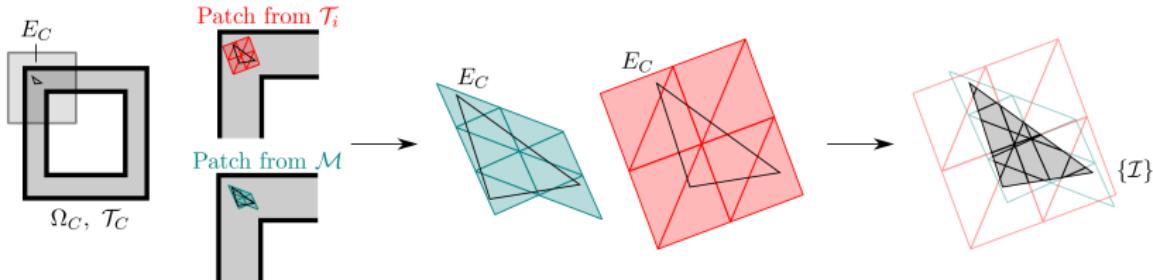
Serial search:

- Test all pairs:
 $O(n^2)$
- Advancing front:
at best $O(n)$

Parallel search: Distribute the work using Ω^C

- 1 Mesh the coupling region Ω^C : \mathcal{T}_C , n_C elems.
→ **No physical system associated!**
- 2 Build “patches” of \mathcal{T}_i and \mathcal{M} overlapping
 $E_C \in \mathcal{T}_C \rightarrow$ **Use advancing front!**
- 3 n_C separate problems: search all intersections
between the patches, *inside* E_C

M. Gander et al., Domain Decomposition Methods in Science and Engineering XVIII, 2009

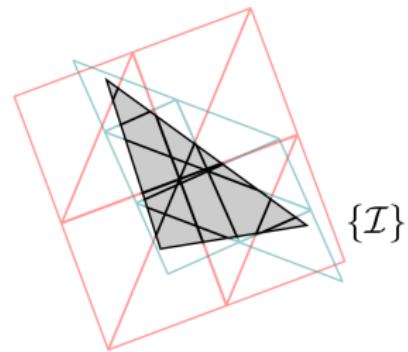


Intersection search: balancing over a computer cluster

- Library used to build intersections: CGAL

<http://www.cgal.org>

- Exact geometry, no edge cases
- Expensive to construct intersection polyhedrons between elements
- Building the set $\{\mathcal{I}\}$: **Expensive**
Find nb. of intersections: **Cheap!**



Balancing:

- 1 Do preliminary, non-balanced intersection search
- 2 Partition \mathcal{T}_C using number of intersections as weight

Intersection search: strong scaling

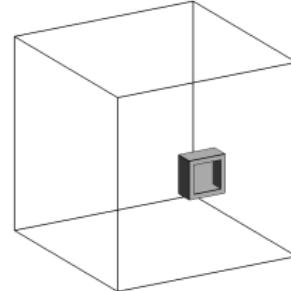
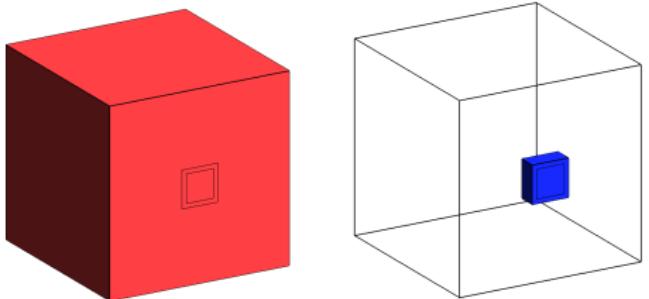
Elem.

$$\mathcal{T}_1 : 1.75 \cdot 10^6$$

$$\mathcal{T}_2 : 4.38 \cdot 10^5$$

$$\mathcal{T}_C : 2.73 \cdot 10^3$$

$$\mathcal{T}_I : 6.66 \cdot 10^6$$



Intersection search: strong scaling

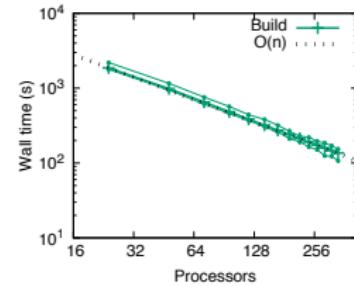
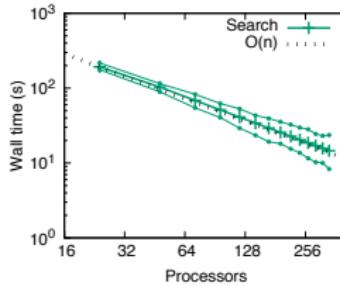
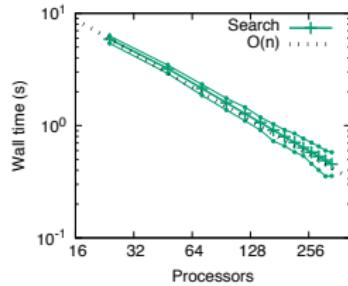


Figure: Search: test all pairs

Xeon E5-2670v3

Elem.

$$\mathcal{T}_1 : 1.75 \cdot 10^6$$

$$\mathcal{T}_2 : 4.38 \cdot 10^5$$

$$\mathcal{T}_C : 2.73 \cdot 10^3$$

$$\mathcal{T}_I : 6.66 \cdot 10^6$$

Figure: Search: advancing front

Figure: Build

- Search time: $\sim 200x$ faster than construction (test all pairs)
- Nb. of elements of \mathcal{T}_C : $\sim 10\times$ nb. of processors
- Advancing front search
 - Needs a construction step \rightarrow Slower

Intersection search: effects of repartitioning

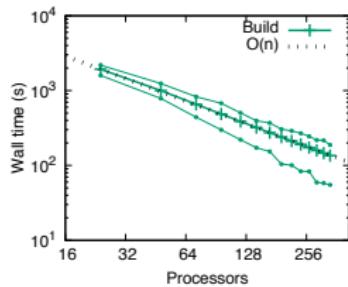


Figure: Build: no repartition

Xeon E5-2670v3

Elem.

$$\mathcal{T}_1 : 1.75 \cdot 10^6$$

$$\mathcal{T}_2 : 4.38 \cdot 10^5$$

$$\mathcal{T}_C : 2.73 \cdot 10^3$$

$$\mathcal{T}_I : 6.66 \cdot 10^6$$

- Repartitioning allows better distribution of intersection construction
- Worth the cost of an extra search step

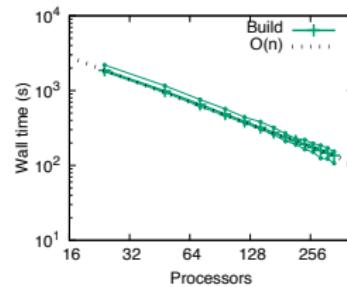


Figure: Build: with repartition

Arlequin coupled solver, CArL

- Main component: `carl::coupled_system` class
- Adds a layer to libMesh's `EquationSystems` → `System` structure

```
std::pair<std::string, libMesh::EquationSystems*> m_macro_EqSys;  
std::map<std::string, libMesh::EquationSystems*> m_micro_EqSys;  
Similar mappings for other systems (intersections, mediator ...) and coupling matrices
```

- Solvers

```
carl::coupled_solver (LATIN, FETI)  
carl::generic_solver_interface (system solvers, can use external programs )
```

- Coupling assemble functions, system parameters ...

CArl: FETI and Arlequin method

- Discretized coupled system:

$$\begin{bmatrix} \mathbf{K}_1 & \mathbf{0} & \mathbf{C}_1^t \\ \mathbf{0} & \mathbf{K}_2 & -\mathbf{C}_2^t \\ \mathbf{C}_1 & -\mathbf{C}_2 & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{u}_1 \\ \mathbf{u}_2 \\ \Phi \end{bmatrix} = \begin{bmatrix} \mathbf{F}_1 \\ \mathbf{F}_2 \\ \mathbf{0} \end{bmatrix}.$$

FETI solver, applied to the Arlequin method

- 1 Find the uncoupled solutions: $\mathbf{K}_i \cdot \mathbf{u}_i^0 = \mathbf{F}_i, i = 1, 2$
- 2 Solve the coupling correction equation:
$$(\mathbf{C}_1 \mathbf{K}_1^{-1} \mathbf{C}_1^t + \mathbf{C}_2 \mathbf{K}_2^{-1} \mathbf{C}_2^t) \cdot \Phi = \mathbf{C}_1 \mathbf{u}_1^0 - \mathbf{C}_2 \mathbf{u}_2^0$$
- 3 Add coupling correction to the solutions:
$$\mathbf{u}_1 = \mathbf{u}_1^0 - \mathbf{K}_1^{-1} \mathbf{C}_1^t \Phi \quad \mathbf{u}_2 = \mathbf{u}_2^0 + \mathbf{K}_2^{-1} \mathbf{C}_2^t \Phi$$

H. B. Dhia et al., Multimodeling of multi-altered structures in the Arlequin framework, European Journal of Computational Mechanics, 2008

C. Farhat et al., A method of finite element tearing and interconnecting and its parallel solution algorithm, International Journal for Numerical Methods in Engineering, 1991

CArl: FETI / PCG solver

$$\underbrace{(\mathbf{C}_1 \mathbf{K}_1^{-1} \mathbf{C}_1^t + \mathbf{C}_2 \mathbf{K}_2^{-1} \mathbf{C}_2^t)}_A \cdot \Phi = \underbrace{\mathbf{C}_1 \mathbf{u}_1^0 - \mathbf{C}_2 \mathbf{u}_2^0}_b$$

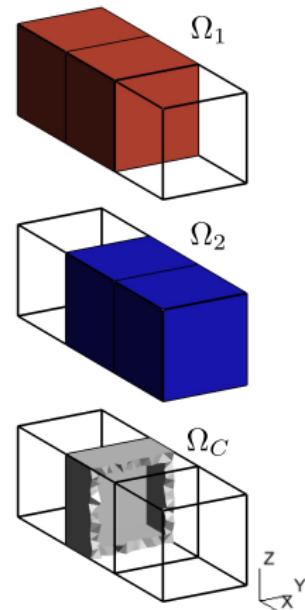
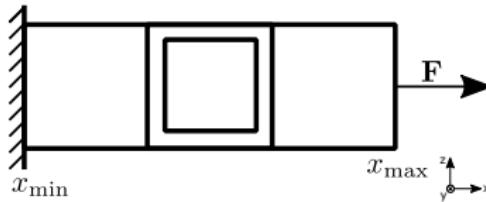
- Avoid explicit construction of \mathbf{A} : Use an iterative solver (CG), and apply \mathbf{A} “by parts”, calling the uncoupled system solvers
- Caveat: one of the models might not have Dirichlet boundary conditions → **ill-conditioned!**
 - 1 Use a **projected** CG algorithm (PCG)
 - 2 Rigid body modes correction: $\mathbf{u}_2 = \mathbf{u}_2^0 + \mathbf{K}_2^+ \mathbf{C}_2^t \Phi + \mathbf{R}_2 \cdot \alpha$
- Possibly time-consuming components of the FETI / PCG solver:
 - Calls to model solvers
 - Preconditioner
 - Rigid body modes projectors (if a model has no boundary conditions)

C. Farhat *et al.*, A method of finite element tearing and interconnecting and its parallel solution algorithm, International Journal for

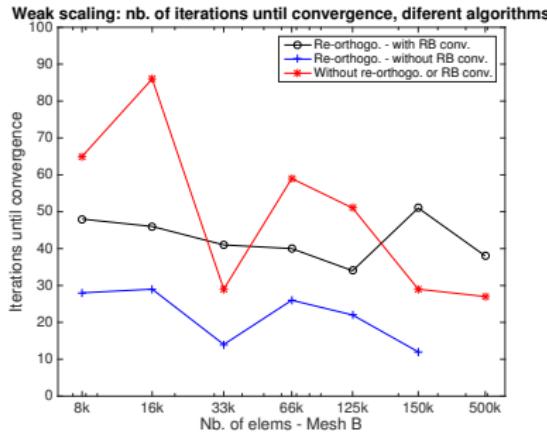
Numerical Methods in Engineering, 1991

CArl: FETI / PCG solver

- Weak scaling of the coupled solver?
- Cost of each part?
- Tests done with traction test
 - **Macro system:** domain Ω_1
 - Clamped x_{\min} face
 - Mediator mesh extracted from this system
 - **Micro system:** domain Ω_2
 - Force applied on x_{\max} face
 - Homogeneous, linear elasticity



CArl / FETI: weak scaling



Elem.

$$\mathcal{T}_1 : 5.0 \cdot 10^5 \quad \mathcal{M} : 1.4 \cdot 10^5$$

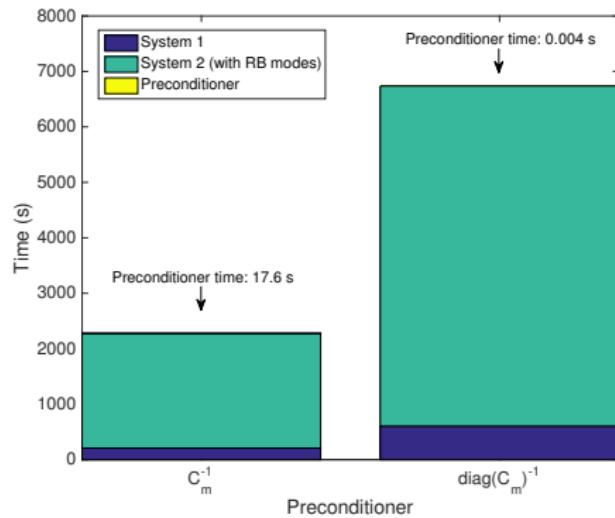
$$\mathcal{T}_2 : 8.0 \cdot 10^4 \sim 5.0 \cdot 10^5$$

(6 ~ 384 procs.)

- Simulation time (and scaling) depends on choice of model solvers
- Weak scaling of only the FETI solver: **iterations until convergence**
- Meshes \mathcal{T}_1 and \mathcal{M} kept unchanged
- Mesh \mathcal{T}_2 : increase number of elems.
- Initial algorithm: erratic number of iterations
- Good scaling when checking convergence of $\mathbf{R}_2 \cdot \alpha$

$$\mathbf{u}_2 = \mathbf{u}_2^0 + \mathbf{K}_2^+ \mathbf{C}_2^t \boldsymbol{\Phi} + \mathbf{R}_2 \cdot \alpha$$

CArl / FETI: computational costs



384 procs. (Xeon E5-2670v3 @ 2.30 GHz)

Elem.

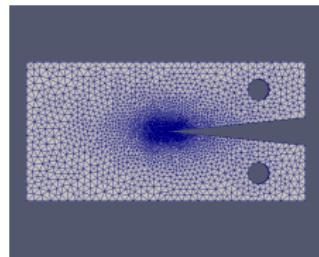
$$\mathcal{T}_1, \mathcal{T}_2 : 1 \cdot 10^6 \quad \mathcal{M} : 2.7 \cdot 10^5$$

- Solve time dominated by calls to model solvers
 - Sys. 1: ~9%
 - Sys. 2: ~90%
 - Projection: ~0.006%
- Preconditioner based on coupling matrix \mathbf{C}_m
- Restricted to mediator space
- \mathbf{C}_m^{-1} : 46 iter. , ~0.7%
- $\text{diag}(\mathbf{C}_m)^{-1}$: 221 iter. , ~0%

Example: linear heterogeneous anisotropic material

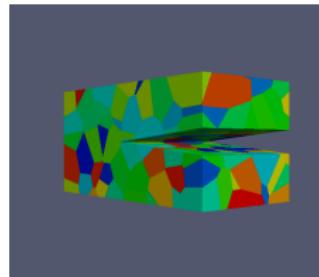
■ Macro system:

- Homogeneous, linear elasticity
- $E = 200 \text{ GPa}$, $\mu = 80 \text{ GPa}$
- Fissure on the xz plan
- z displacements enforced on the holes
- Clamped left surface

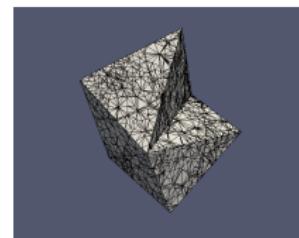
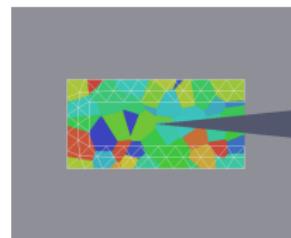
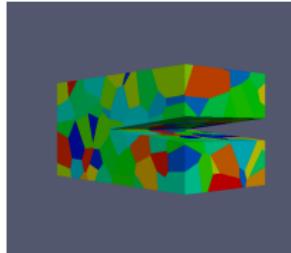
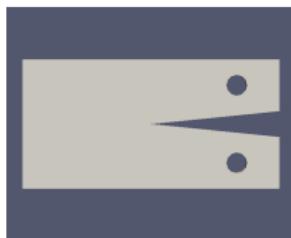


■ Micro system:

- Heterogeneous, anisotropic elasticity
- $c_{11} = 198 \text{ GPa}$, $c_{12} = 125 \text{ GPa}$, $c_{44} = 122 \text{ GPa}$
- Elasticity tensor: cubic symmetry, random grain orientations (isotropic distribution)
- No direct boundary conditions / forces
- Mesh generated from Voronoi diagram
- 250 grains



Example: linear heterogeneous anisotropic material



Meshes

Macro (homogeneous):

Micro (250 grains, anisotropic):

Coupling region:

Intersection mesh:

Elem.

$3.4 \cdot 10^4$

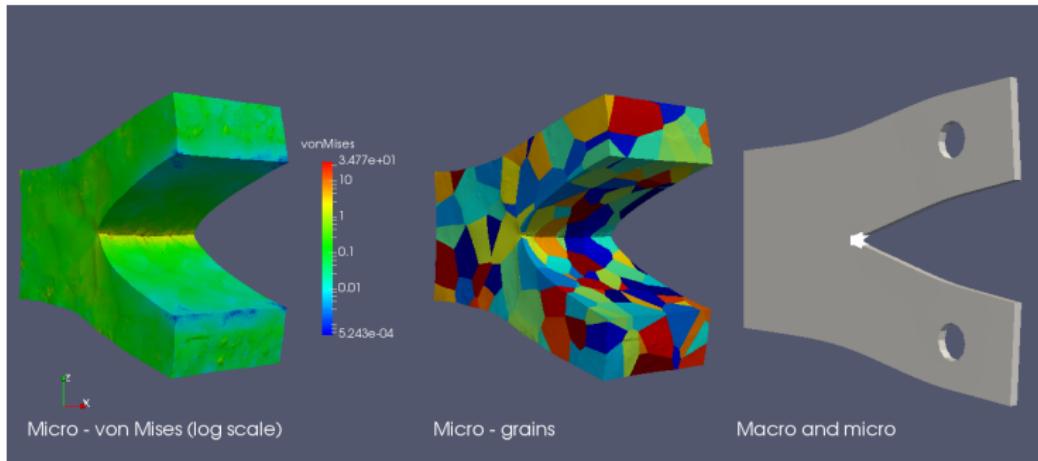
$1.1 \cdot 10^6$

$1.1 \cdot 10^3$

$7.2 \cdot 10^6$

- Macro mesh element length: $\sim 10x$ smaller near fissure
- Micro mesh element length: $\sim 10x$ smaller than the smallest macro system element

Example: linear heterogeneous anisotropic material



Coupling test

- 96 processors (Xeon E5-2670v3 @ 2.30 GHz), wall time
 - Intersections: ~ 11 minutes Solve: ~ 1 hour

Conclusion

- Construction of a scalable algorithm to find the intersections between overlapping, non-conformal meshes
 - No need to adapt the models' meshes to build the coupling
 - No need to add construction constraints to the random meshes
- Development of a scalable, parallel coupled system solver

Current work

- Applying the current code to the homogenization problem
 - **Macro:** homogeneous, deterministic
 - **Micro:** stochastic, non-linear, anisotropic polycrystal

T. M. Schlittler, R. Cottreau, Fully scalable implementation of a volume coupling scheme for the modeling of polycrystalline materials, Computational Mechanics (article being finalized)

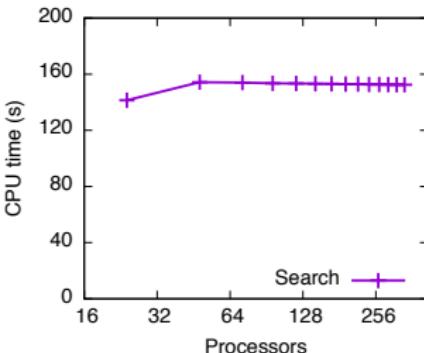


Github

<https://github.com/cottreau/CArI>

Thanks for your attention!

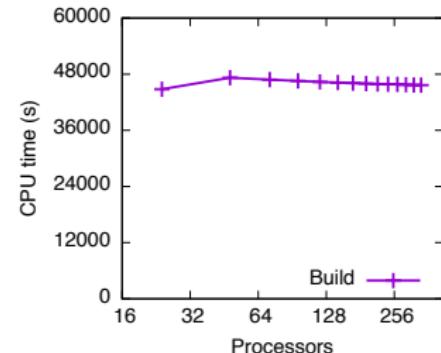
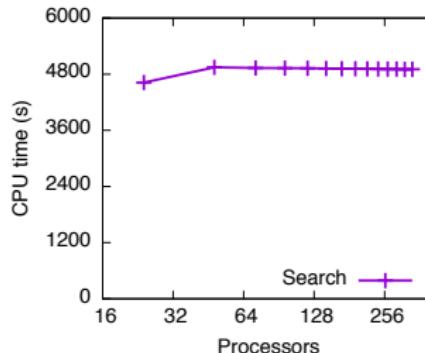
Intersection search: scaling



Xeon E5-2695V2
@ 2.4GHz

Elem.

$$\begin{aligned} T_1 &: 1.75 \cdot 10^6 \\ T_2 &: 4.38 \cdot 10^5 \\ T_C &: 2.73 \cdot 10^3 \\ T_I &: 6.66 \cdot 10^6 \end{aligned}$$



- Search time: $\sim 200\times$ faster than construction (test all pairs)
- Nb. of elements of T_C : $\sim 10\times$ nb. of processors
- Advancing front search
 - Needs a construction step \rightarrow Slower